


ECSE-626

Statistical Computer Vision

Gaussian Mixture Models, EM, Kernel Density Estimation

Introduction



- We often want to estimate a probability density function from a set of observations.



?



Modeling Densities

- Given a set of observations, how do we model a probability density?



Modeling Densities

- Given a set of observations, how do we model a probability density?



- Often we model a set of samples as having a Gaussian density

Gaussian Densities



- Gaussian densities can be compactly represented by their mean and variance
- Sample mean and variance can be computed in linear time

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$

$$s^2 = \frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})^2$$

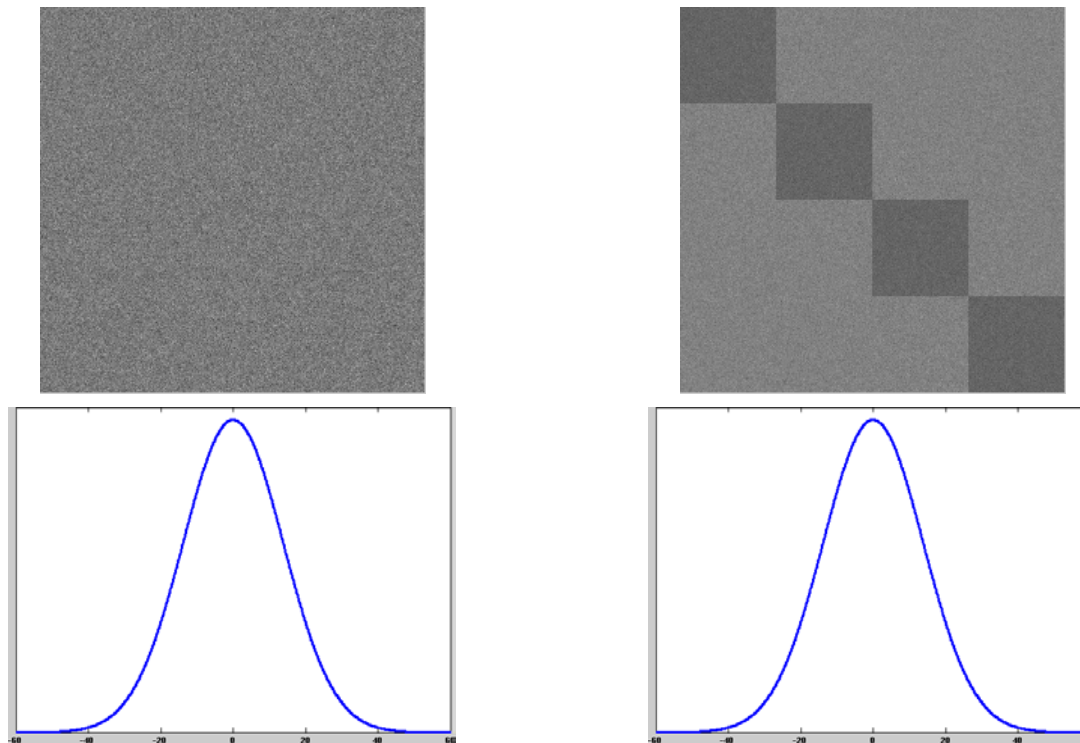
Gaussian Densities

- Gaussian densities have many convenient mathematical properties and serve as a reasonable approximation for many true densities.



Gaussian Densities

- However, many real densities are not well modeled by Gaussians.





Gaussian Densities

- Parametric densities such as Gaussians, and all exponential densities, are unimodal (have a single local maximum).
- Many practical problems involve multi-modal densities.



Multi-modal Densities

- Gaussian is poor representation of densities that have multiple modes.





Gaussian Mixture Model (GMM)

- Allow us to represent densities as weighted sum (or mixture) of multiple Gaussians



Gaussian Mixture Model (GMM)

- Allow us to represent densities as weighted sum (or mixture) of multiple Gaussians



Gaussian Mixture Models (GMM)

- Mixtures of Gaussian functions are well-suited to modeling clusters of points.
- Each cluster is assigned a Gaussian, with its mean in the middle of the cluster and with a standard deviation measuring its spread.

Gaussian Mixture Models (GMM)

K : number of gaussian components

μ_k : mean of k^{th} gaussian component

σ_k : std of k^{th} gaussian component



α_k : mixing probability (weight) of k^{th} gaussian component

$$\Theta = \{\alpha_1, \dots, \alpha_k, \mu_1, \dots, \mu_k, \sigma_1, \dots, \sigma_k\}$$

$$f(x|\Theta) = \sum_{k=1}^K \alpha_k \mathcal{N}(x; \mu_k, \sigma_k)$$



Gaussian Mixture Models

$$K = 2, \mu = \{-21, 7.00\}, \sigma = \{30, 50\}, \alpha = \{0.25, 0.75\}$$



Gaussian Mixture Models


$$K = 2, \mu = \{-21, 7.00\}, \sigma = \{30, 50\}, \alpha = \{0.25, 0.75\}$$




$k = 1$

Gaussian Mixture Models

$$K = 2, \mu = \{-21, 7.00\}, \sigma = \{30, 50\}, \alpha = \{0.25, 0.75\}$$





$k = 2$





Gaussian Mixture Models

$$K = 2, \mu = \{-21, 7.00\}, \sigma = \{30, 50\}, \alpha = \{0.25, 0.75\}$$





Gaussian Mixture Models

$$K = 2, \mu = \{-21, 7.00\}, \sigma = \{30, 50\}, \alpha = \{0.25, 0.75\}$$




Gaussian Mixture Models

$$K = 2, \mu = \{-21, 7.00\}, \sigma = \{30, 50\}, \alpha = \{0.25, 0.75\}$$




Gaussian Mixture Models

$$K = 5, \mu = \{-5, 0, 3, 5, 8\}, \sigma = \{1, 2, 1, 2, 1\}, \alpha = \{0.2, 0.2, 0.4, 0.1, 0.1\}$$




Gaussian Mixture Models

$$K = 5, \mu = \{-5, 0, 3, 5, 8\}, \sigma = \{1, 2, 1, 2, 1\}, \alpha = \{0.2, 0.2, 0.4, 0.1, 0.1\}$$



Gaussian Mixture Models




$$K = 5, \mu = \{-5, 0, 3, 5, 8\}, \sigma = \{1, 2, 1, 2, 1\}, \alpha = \{0.2, 0.2, 0.4, 0.1, 0.1\}$$



Gaussian Mixture Models

- For most real problems, we don't know how many or which Gaussians the data are sampled from.
- How do we automatically determine the parameters of a GMM from observed samples?

Unsupervised Learning - GMMs



- Given an image we want to determine parameters Θ :

$$\Theta = \{\alpha_1, \dots, \alpha_k, \mu_1, \dots, \mu_k, \sigma_1, \dots, \sigma_k\}$$

Maximum Likelihood

- Assume we know form of model (GMM) :

$$p(x|\Theta) = \sum_{k=1}^K \alpha_k \mathcal{N}(x; \mu_k, \sigma_k)$$

Maximum Likelihood

- Assume we know form of model (GMM) :


$$p(x|\Theta) = \sum_{k=1}^K \alpha_k \mathcal{N}(x; \mu_k, \sigma_k)$$

- Assume observed samples, x_i , are i.i.d.:

$$p(\mathcal{X}|\Theta) = \prod_{i=1}^N p(x_i|\Theta)$$

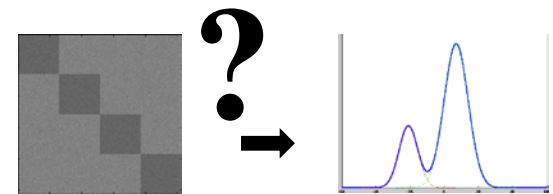
Maximum Likelihood

$$p(\mathcal{X}|\Theta) = \prod_{i=1}^N p(x_i|\Theta)$$



N : # samples, i : sample (pixel) index, Θ : GMM parameters, K : # components in GMM

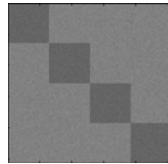

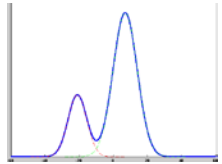
Maximum Likelihood

$$p(\mathcal{X}|\Theta) = \prod_{i=1}^N p(x_i|\Theta)$$


$$= \prod_{i=1}^N \sum_{k=1}^K \alpha_k \mathcal{N}(x_i; \mu_k, \sigma_k)$$

N : # samples, i : sample (pixel) index, Θ : GMM parameters, K : # components in GMM

Maximum Likelihood

$$p(\mathcal{X}|\Theta) = \prod_{i=1}^N p(x_i|\Theta)$$





$$= \prod_{i=1}^N \sum_{k=1}^K \alpha_k \mathcal{N}(x_i; \mu_k, \sigma_k)$$

- Want to find model parameters, Θ , that maximize likelihood.

N : # samples, i : sample (pixel) index, Θ : GMM parameters, K : # components in GMM

Maximum Likelihood for GMM


$$\Theta^* = \operatorname{argmax}_{\Theta} p(\mathcal{X}|\Theta)$$



Maximum Likelihood for GMM


$$\Theta^* = \operatorname{argmax}_{\Theta} p(\mathcal{X}|\Theta)$$

$$= \operatorname{argmax}_{\Theta} \log p(\mathcal{X}|\Theta)$$



Maximum Likelihood for GMM

$$\Theta^* = \operatorname{argmax}_{\Theta} p(\mathcal{X}|\Theta)$$




$$= \operatorname{argmax}_{\Theta} \log p(\mathcal{X}|\Theta)$$

$$= \operatorname{argmax}_{\Theta} \log \prod_{i=1}^N \sum_{k=1}^K \alpha_k \mathcal{N}(x_i; \mu_k, \sigma_k)$$

Maximum Likelihood for GMM

$$\Theta^* = \operatorname{argmax}_{\Theta} p(\mathcal{X}|\Theta)$$



$$= \operatorname{argmax}_{\Theta} \log p(\mathcal{X}|\Theta)$$

$$= \operatorname{argmax}_{\Theta} \log \prod_{i=1}^N \sum_{k=1}^K \alpha_k \mathcal{N}(x_i; \mu_k, \sigma_k)$$

$$= \operatorname{argmax}_{\Theta} \sum_{i=1}^N \log \left(\sum_{k=1}^K \alpha_k \mathcal{N}(x_i; \mu_k, \sigma_k) \right)$$

Maximum Likelihood for GMM

$$\Theta^* = \underset{\Theta}{\operatorname{argmax}} \sum_{i=1}^N \log \left(\sum_{k=1}^K \alpha_k \mathcal{N}(x_i; \mu_k, \sigma_k) \right)$$

- Difficult to optimize because contains log of summation



Expectation Maximization (EM)

- Iterative algorithm for finding maximum-likelihood estimate of model parameters from a given data set.
- **Expectation** : determine which Gaussian component(s) each sample comes from based on current estimate of model parameters
- **Maximization** : maximize model parameters given current estimates of Gaussian sources of each sample

GMMs for Segmentation


- We have started to look at the use of GMMs / EM for estimating multi-modal probability densities.
- GMMs / EM are also often used for segmentation or clustering tasks
- EM can be thought of as conceptually similar to K-means
- Recall K-means for segmentation....

Image Segmentation: toy example





- These intensities define the three groups
- We could label every pixel in the image according to which of these primary intensities it is
 - i.e., *segment* the image based on the intensity feature.
- What if the image isn't quite so simple?

Image Segmentation: toy example



input image



input image






Image Segmentation: toy example




input image




- Now how to determine the three main intensities that define our groups?
- We need to *cluster*

Image Segmentation: Clustering



- Goal: choose three *centers* as the representative intensities, and label every pixel according to which of these centers it is nearest to

Image Segmentation: Clustering




- Best cluster centers are those that minimize SSD between all points and their nearest cluster center c_i


$$\sum_{\text{cluster } i} \sum_{\text{points } p \text{ in cluster } i} \|p - c_i\|^2$$

Image Segmentation: Clustering


- With this objective, it is a “chicken and egg” problem
 - If we knew the **cluster centers**, we could allocate points to groups by assigning each to its closest center



- If we knew the group memberships, we could get the centers by computing the mean per group




K-means Clustering Algorithm

- Randomly initialize the K cluster centers, and iterate between the two steps we just saw
 1. Randomly initialize the cluster centers c_1, \dots, c_K
 2. Given the cluster centers, determine points in each cluster
 - For each point p , find the closest c_i . Put p into cluster i
 3. Given points in each cluster, solve for c_i
 - Set c_i to be the mean of points in cluster i
 4. If c_i have changed, repeat Step 2
- 

K-means Clustering

- Converging to a local minimum







GMMs for Segmentation

- GMMs / EM are also often used for segmentation or clustering tasks
- We can think of each of the K components in our GMM as representing a distinct *class* and each sample as having a membership to one or more of the K classes.


GMMs for Segmentation

- We can model our image as a GMM consisting of K components, each representing a distinct class.
- We can then assign a likelihood of each of the K classes for every sample in our image

GMMs for Segmentation




$$p(y_i = 1 | x_i, \Theta)$$




$$p(y_i = 2 | x_i, \Theta)$$

Expectation Maximization



- If we know which class (gaussian component) each sample came from, then we can estimate the parameters of our GMM.
- If we know the parameters of the GMM, we can determine the likelihood of each sample belonging to each class.

Expectation Maximization



- In practice we don't know the GMM parameters or the classes of each sample but only have the image.

Expectation Maximization

$$\Theta^* = \underset{\Theta}{\operatorname{argmax}} \sum_{i=1}^N \log \left(\sum_{k=1}^K \alpha_k \mathcal{N}(x_i; \mu_k, \sigma_k) \right)$$

- To simplify our optimization, we introduce a hidden parameter, Y , that acts as a class label (corresponding to one of K gaussian clusters) for each sample, x_i :

$$\mathcal{Y} = \{y_i\}_{i=1}^N, y_i \in 1, \dots, K$$

Expectation Maximization

$$\mathcal{Y} = \{y_i\}_{i=1}^N, y_i \in 1, \dots, K$$

- For segmentation, we can think of each y_i as a class label for sample x_i

Expectation Maximization

$$\mathcal{Y} = \{y_i\}_{i=1}^N, y_i \in 1, \dots, K$$

- For segmentation, we can think of each y_i as a class label for sample x_i
- For density estimation, y_i acts as a hidden parameter that will allow us to solve for our GMM parameters, Θ .

Expectation Maximization

$$\mathcal{Y} = \{y_i\}_{i=1}^N, y_i \in 1, \dots, K$$

- For segmentation, we can think of each y_i as a class label for sample x_i
- For density estimation, y_i acts as a hidden parameter that will allow us to solve for our GMM parameters, Θ .
- In general, each y_i will be represented as a distribution over all K possible classes.

Expectation Maximization


$$p(x_i, y_i | \Theta) = p(x_i | y_i, \Theta) p(y_i | \Theta)$$

Expectation Maximization

$$p(x_i, y_i | \Theta) = p(x_i | y_i, \Theta) p(y_i | \Theta)$$

- We can simplify by observing the following:

$$p(y_i | \Theta) = \alpha_{y_i}$$




Expectation Maximization

$$p(x_i, y_i | \Theta) = p(x_i | y_i, \Theta) p(y_i | \Theta)$$

- We can simplify by observing the following:

$$p(y_i | \Theta) = \alpha_{y_i}$$

$$p(x_i | y_i, \Theta) = \mathcal{N}(x_i; \mu_{y_i}, \sigma_{y_i})$$



Expectation Maximization

$$p(x_i, y_i | \Theta) = p(x_i | y_i, \Theta) p(y_i | \Theta)$$

- We can simplify by observing the following:

$$p(y_i | \Theta) = \alpha_{y_i}$$

$$p(x_i | y_i, \Theta) = \mathcal{N}(x_i; \mu_{y_i}, \sigma_{y_i})$$

$$p(x_i, y_i | \Theta) = \mathcal{N}(x_i; \mu_{y_i}, \sigma_{y_i}) \alpha_{y_i}$$

Expectation Maximization

$$p(\mathcal{X}, \mathcal{Y} | \Theta) = \prod_{i=1}^N p(x_i, y_i | \Theta)$$

Expectation Maximization

$$\begin{aligned} p(\mathcal{X}, \mathcal{Y} | \Theta) &= \prod_{i=1}^N p(x_i, y_i | \Theta) \\ &= \prod_{i=1}^N \mathcal{N}(x_i; \mu_{y_i}, \sigma_{y_i}) \alpha_{y_i} \end{aligned}$$

Expectation Maximization

$$\Theta^* = \operatorname{argmax}_{\Theta} p(\mathcal{X}, \mathcal{Y} | \Theta)$$

Expectation Maximization

$$\begin{aligned}\Theta^* &= \operatorname{argmax}_{\Theta} p(\mathcal{X}, \mathcal{Y} | \Theta) \\ &= \operatorname{argmax}_{\Theta} \log p(\mathcal{X}, \mathcal{Y} | \Theta)\end{aligned}$$

Expectation Maximization

$$\begin{aligned}\Theta^* &= \operatorname{argmax}_{\Theta} p(\mathcal{X}, \mathcal{Y} | \Theta) \\ &= \operatorname{argmax}_{\Theta} \log p(\mathcal{X}, \mathcal{Y} | \Theta) \\ &= \operatorname{argmax}_{\Theta} \log \prod_{i=1}^N \mathcal{N}(x_i; \mu_{y_i}, \sigma_{y_i}) \alpha_{y_i}\end{aligned}$$

Expectation Maximization

$$\begin{aligned}\Theta^* &= \operatorname{argmax}_{\Theta} p(\mathcal{X}, \mathcal{Y} | \Theta) \\ &= \operatorname{argmax}_{\Theta} \log p(\mathcal{X}, \mathcal{Y} | \Theta) \\ &= \operatorname{argmax}_{\Theta} \log \prod_{i=1}^N \mathcal{N}(x_i; \mu_{y_i}, \sigma_{y_i}) \alpha_{y_i} \\ &= \operatorname{argmax}_{\Theta} \sum_{i=1}^N \log \alpha_{y_i} \mathcal{N}(x_i; \mu_{y_i}, \sigma_{y_i})\end{aligned}$$


Expectation Maximization

$$\Theta^* = \operatorname{argmax}_{\Theta} \sum_{i=1}^N \log \alpha_{y_i} \mathcal{N}(x_i; \mu_{y_i}, \sigma_{y_i})$$

- One problem : we don't actually have y_i ...
- Iterative optimization
 - E-step : estimate y_i based on current estimate of parameters
 - M-step : find parameters, Θ , that maximize expectation of log-likelihood
- Requires initial estimate of parameters, Θ , or of y_i

E-step


- Determine density of the hidden parameters, y_i , based on current estimate of model parameters



$$p(y_i | x_i, \Theta^t)$$


E-step

$$p(y_i|x_i, \Theta^t) = \frac{p(x_i|y_i, \Theta^t)p(y_i|\Theta^t)}{p(x_i|\Theta^t)}$$




E-step

$$\begin{aligned} p(y_i|x_i, \Theta^t) &= \frac{p(x_i|y_i, \Theta^t)p(y_i|\Theta^t)}{p(x_i|\Theta^t)} \\ &= \frac{\mathcal{N}(x_i; \mu_{y_i}^t, \sigma_{y_i}^t) \alpha_{y_i}^t}{p(x_i|\Theta^t)} \end{aligned}$$




E-step

$$\begin{aligned}
 p(y_i|x_i, \Theta^t) &= \frac{p(x_i|y_i, \Theta^t)p(y_i|\Theta^t)}{p(x_i|\Theta^t)} \\
 &= \frac{\mathcal{N}(x_i; \mu_{y_i}^t, \sigma_{y_i}^t) \alpha_{y_i}^t}{p(x_i|\Theta^t)} \\
 &= \frac{\mathcal{N}(x_i; \mu_{y_i}^t, \sigma_{y_i}^t) \alpha_{y_i}^t}{\sum_{k=1}^K \alpha_k^t \mathcal{N}(x_i; \mu_k^t, \sigma_k^t)}
 \end{aligned}$$



M-step

- Find model parameters that maximize expectation of log likelihood based on current estimate of density of hidden parameters




M-step

$$\mathbb{E}[\log p(\mathcal{X}, \mathcal{Y} | \Theta) | \mathcal{X}, \Theta^t] = \sum_{\mathbf{y} \in \mathcal{Y}} \log p(\mathcal{X}, \mathcal{Y} | \Theta) p(\mathbf{y} | \mathcal{X}, \Theta^t)$$

$$Q(\Theta, \Theta^t) = \mathbb{E}[\log p(\mathcal{X}, \mathcal{Y} | \Theta) | \mathcal{X}, \Theta^t]$$

$$\Theta^{t+1} = \operatorname{argmax}_{\Theta} Q(\Theta, \Theta^t)$$



M-step

- After a long derivation, we can get analytical equations for our M-step if using GMMs:

$$\alpha_k^{t+1} = \frac{1}{N} \sum_{i=1}^N p(y_i = k | x_i, \Theta^t)$$

$$\mu_k^{t+1} = \frac{\sum_{i=1}^N x_i p(y_i = k | x_i, \Theta^t)}{\sum_{i=1}^N p(y_i = k | x_i, \Theta^t)}$$

$$\sigma_k^{t+1} = \frac{\sum_{i=1}^N p(y_i = k | x_i, \Theta^t) (x_i - \mu_k^t)^2}{\sum_{i=1}^N p(y_i = k | x_i, \Theta^t)}$$

M-step

- The update equations are actually quite intuitive:

$$\alpha_k^{t+1} = \frac{1}{N} \sum_{i=1}^N p(y_i = k | x_i, \Theta^t)$$

- This represents the total relative weight of samples that are of class k (or gaussian component k), based on estimates of class labels (or hidden parameters), y_i

M-step

$$\mu_k^{t+1} = \frac{\sum_{i=1}^N x_i p(y_i = k | x_i, \Theta^t)}{\sum_{i=1}^N p(y_i = k | x_i, \Theta^t)}$$


- This represents the weighted sample mean of samples that are of class k (or gaussian component k), based on estimates of class labels (or hidden parameters), y_i

M-step

$$\sigma_k^{t+1} = \frac{\sum_{i=1}^N p(y_i = k | x_i, \Theta^t) (x_i - \mu_k^t)^2}{\sum_{i=1}^N p(y_i = k | x_i, \Theta^t)}$$


- This represents the weighted sample variance of samples that are of class k (or gaussian component k), based on estimates of class labels (or hidden parameters), y_i

EM in practice






- Assume $K = 2$
- Initialize y_i such that samples with lowest 50% of intensities are assigned a label of 1, while highest 50% are assigned to label 2

EM in practice





- Assume $K = 2$
- Initialize y_i such that samples with lowest 50% of intensities are assigned a label of 1, while highest 50% are assigned to label 2


EM in practice : M-step



$$p(\mathcal{X}|\Theta)$$





$$p(y_i = 1|x_i, \Theta)$$





$$p(y_i = 2|x_i, \Theta)$$


EM in practice : E-step



$$p(\mathcal{X}|\Theta)$$





$$p(y_i = 1|x_i, \Theta)$$





$$p(y_i = 2|x_i, \Theta)$$


EM in practice : M-step



$$p(\mathcal{X}|\Theta)$$





$$p(y_i = 1|x_i, \Theta)$$





$$p(y_i = 2|x_i, \Theta)$$


EM in practice : E-step



$$p(\mathcal{X}|\Theta)$$





$$p(y_i = 1|x_i, \Theta)$$





$$p(y_i = 2|x_i, \Theta)$$


EM in practice : M-step



$$p(\mathcal{X}|\Theta)$$





$$p(y_i = 1|x_i, \Theta)$$





$$p(y_i = 2|x_i, \Theta)$$


EM in practice : E-step



$$p(\mathcal{X}|\Theta)$$





$$p(y_i = 1|x_i, \Theta)$$





$$p(y_i = 2|x_i, \Theta)$$


EM in practice : M-step



$$p(\mathcal{X}|\Theta)$$





$$p(y_i = 1|x_i, \Theta)$$





$$p(y_i = 2|x_i, \Theta)$$


EM in practice



$$p(\mathcal{X}|\Theta)$$



$$p(y_i = 1|x_i, \Theta)$$




$$p(y_i = 2|x_i, \Theta)$$

EM in practice




- Want to segment brain MRI into 3 tissue classes:
 - cerebro-spinal fluid (csf)
 - gray matter
 - white matter

EM in practice





- Initialize y_i by assigning lowest 1/3 of intensities to csf, middle 1/3 to gray matter, and highest 1/3 to white matter

EM in practice



- Initialize y_i by assigning lowest 1/3 of intensities to csf, middle 1/3 to gray matter, and highest 1/3 to white matter

EM in practice : M-step




$$p(\mathcal{X}|\Theta)$$



$$p(y_i = 1|x_i, \Theta)$$




$$p(y_i = 2|x_i, \Theta)$$




$$p(y_i = 3|x_i, \Theta)$$

EM in practice : E-step



$$p(\mathcal{X}|\Theta)$$



$$p(y_i = 1|x_i, \Theta)$$




$$p(y_i = 2|x_i, \Theta)$$





$$p(y_i = 3|x_i, \Theta)$$

EM in practice : M-step



$$p(\mathcal{X}|\Theta)$$



$$p(y_i = 1|x_i, \Theta)$$




$$p(y_i = 2|x_i, \Theta)$$




$$p(y_i = 3|x_i, \Theta)$$

EM in practice : E-step



$$p(\mathcal{X}|\Theta)$$



$$p(y_i = 1|x_i, \Theta)$$





$$p(y_i = 2|x_i, \Theta)$$




$$p(y_i = 3|x_i, \Theta)$$

EM in practice : M-step



$$p(\mathcal{X}|\Theta)$$



$$p(y_i = 1|x_i, \Theta)$$





$$p(y_i = 2|x_i, \Theta)$$




$$p(y_i = 3|x_i, \Theta)$$

EM in practice



$$p(\mathcal{X}|\Theta)$$



$$p(y_i = 1|x_i, \Theta)$$



$$p(y_i = 2|x_i, \Theta)$$




$$p(y_i = 3|x_i, \Theta)$$

EM in practice





- What if we start with a random initialization?

EM in practice





- What if we start with a random initialization?


EM in practice




$$p(\mathcal{X}|\Theta)$$



$$p(y_i = 1|x_i, \Theta)$$




$$p(y_i = 2|x_i, \Theta)$$




$$p(y_i = 3|x_i, \Theta)$$

EM in practice





- Assume $K = 2$
- Random initialization

EM in practice





- Assume $K = 2$
- Random initialization


EM in practice



$$p(\mathcal{X}|\Theta)$$




$$p(y_i = 1|x_i, \Theta)$$




$$p(y_i = 2|x_i, \Theta)$$

EM in practice







- What if we set $K = 4$?
- Initialize by setting lowest 25% of intensities to class 1, etc...

EM in practice



- What if we set $K = 4$?
- Initialize by setting lowest 25% of intensities to class 1, etc...

EM in practice



$$p(y_i = 1|x_i, \Theta)$$



$$p(y_i = 2|x_i, \Theta)$$




$$p(y_i = 3|x_i, \Theta)$$




$$p(y_i = 4|x_i, \Theta)$$

EM in practice










- What if we set $K = 4$ and initialize randomly?

EM in practice



- What if we set $K = 4$ and initialize randomly?

EM in practice




$$p(y_i = 1|x_i, \Theta)$$

$$p(y_i = 2|x_i, \Theta)$$


$$p(y_i = 3|x_i, \Theta)$$

$$p(y_i = 4|x_i, \Theta)$$


Effect of K and initialization




$K = 3$



$K = 3$, random



$K = 4$



$K = 4$, random

Expectation Maximization

- Iterate until convergence – guaranteed to converge to at least a local maximum
- More details on the derivation of the update equations (and EM in general) can be found here:

<http://www.icsi.berkeley.edu/ftp/global/pub/techreports/1997/tr-97-021.pdf>

GMMs / Expectation Maximization


- GMM parameters provide multimodal estimate of a probability density from a set of samples
- Hidden parameters, y_i , can be treated as class labels if used in a segmentation or clustering framework and provide likelihood of class k at each sample, where class k is parametrized by the k^{th} Gaussian component of our GMM

GMMs

- Drawbacks
 - Need to know K or have some way to determine K
 - Final distribution depends on initialization of model parameters
 - Cannot approximate all arbitrary distributions with limited number of Gaussians
 - May take long time to converge
- Advantages
 - Parametric : still have relatively compact representation
 - Can approximate many multimodal distributions reasonably well with relatively small number of Gaussians
 - Relatively easy to implement


Non-parametric densities

- Non-parametric representations make no assumptions about the form of underlying density



Histogram

- Simple non-parametric representation
- Bins too large : loss of resolution
- Bins too small : zero probability “holes” in density



Kernel Density Estimation

- Method of building non-parametric representations of densities from data samples
- Densities represented as expected value of “smoothed” samples
- Also known as “Parzen Window” method

$$\hat{f}(x) = \frac{1}{N} \sum_{i=1}^N K(x - x_i)$$

Kernel Density Estimation

$$\hat{f}(x) = \frac{1}{N} \sum_{i=1}^N K(x - x_i)$$

- $K(x)$ is kernel (or window) function

$$\int K(x) = 1$$


Kernel Density Estimation

- Gaussian is often used as smoothing kernel


$$\hat{f}(x) = \frac{1}{N} \sum_{i=1}^N \mathcal{N}(\mu = x_i, \sigma)$$

- Density represented as sum of Gaussians, where each Gaussian represents a data sample and has variance σ^2 and mean value equal to the value of the data sample.


Parzen Windows



Parzen Windows




Parzen Windows





Parzen Windows

- Require bandwidth parameter (covariance of Gaussian kernel)
 - Determines “smoothness” of representation
 - Similar to bin width selection
 - Choice of kernel bandwidth generally based on # data samples and variance of data
 - Several methods exist to select bandwidth based on available sample data

Kernel Bandwidth




Parzen Windows




Black = 3 component GMM

Parzen Windows





Black = 3 component GMM
Red = Parzen Window (BW = 1)

Parzen Windows



Black = 3 component GMM
Red = Parzen Window (BW = 4)

Parzen Windows



Black = 3 component GMM
Red = Parzen Window (BW = 16)

Densities using Parzen Windows

- Advantages
 - Can represent arbitrary densities
 - Does not require initialization
 - Methods exist to select appropriate kernel bandwidth based on available data samples
- Disadvantages
 - Can not represent density in a compact form
 - Need to determine kernel bandwidth
 - Non-trivial implementation details, especially for higher-dimensional densities

Summary

- GMMs
 - Can approximate multi-modal densities in a relatively compact form.
 - Given K , EM allows us to generate an estimate of the underlying probability density of our samples.
 - For segmentation or clustering tasks, we can model our samples as a GMM and use EM to determine the underlying cluster from which each sample came.
- Parzen Windows
 - Can represent arbitrary probability densities without any initialization
 - Non-compact representation
 - More difficult to implement